

■ Singularity Health Center

Senior Engineering Manager plan for a hyperscale Agent Platform initiative

Deliver customer-trust outcomes while building a reusable real-time health platform.

The assignment is not only an architecture exercise. It tests whether we can deliver a strategic platform while absorbing live operational load, dependency delays, and senior-engineer disagreement.

■ Customer Outcome

Trusted visibility into agent health across millions of endpoints.

■ Platform Outcome

Reusable telemetry, coalescing, alert lifecycle, and replay capabilities.

■ Execution Outcome

MVP in controlled rollout, GA after scale, correctness, and migration gates.

■ Org Outcome

Team stays focused without ignoring current customer pain.

■ <5 min detection

■ <200 ms API p95

■ Billions events/day

■ Replay + backfill

■ Tenant rollout

Leadership thesis: protect customers now, reduce legacy drag, and land the new platform through staged scope and evidence-based decisions.

02

Scope, SLOs, And Delivery Guardrails

Define what matters before debating technology choices

MVP Signals

Offline / connectivity loss
Agent disabled
Anti-tamper disabled
Low disk / resource risk

GA Expansion

Alert lifecycle
Suppression and drill-down
Shadow migration from heartbeat
Tenant-wide rollout

Explicitly Deferred


Generic rule builder
Advanced analytics
Custom notifications
Full legacy retirement

SLO	Target	Why It Matters
Detection freshness	99% < 5 min	Trustworthy health state
Dashboard latency	p95 < 200 ms	Operator workflow speed
Silent gap detection	< 5 min to page	No stale green dashboards
Accepted event loss	0 after durable bus	Replayable correctness

◆ Guardrail: every MVP feature must improve customer actionability or de-risk the platform launch.

Reference Architecture

Cloud-neutral, replayable, and intentionally separated into hot, read, and cold paths



Hot Path

Event-time detection, coalescing, alert transitions, freshness SLO.

Read Path


Precomputed dashboard views, bounded queries, freshness metadata.

Cold Path

Long-term retention, backfill, analytics, and rule tuning.

■ Read Path And Latency Budget

The UI should query health state, not raw telemetry



■ Budget Segment	Target	Design Choice
Auth / gateway	20–40 ms	Tenant-scoped claims & request shaping
API orchestration	40–60 ms	No fanout over raw telemetry
Read / search	70–90 ms	Precomputed views & bounded filters
Serialization / network	20–30 ms	Slim response models & pagination

■ Graceful Degradation

If search lags, serve current alert state with explicit freshness metadata.

■ Query Discipline

Pagination, capped time windows, no tenant-wide scans from UI paths.



■ Correctness Model

At-least-once ingestion plus idempotent alert transitions.

■ Alert Coalescing Design

Turn noisy telemetry into actionable health state within the 5-minute SLO

Example: 50 anti-tamper disabled events over 10 minutes



■ State Machine

Open → update evidence count → suppress duplicate
→ resolve when healthy signal appears.

■ Auditability

Store raw event references and rule version, not every
raw event in the UI path.

■ Decision


Use stream processing for coalescing; keep rule scope
narrow for MVP.

- Why not cron-first: late detection, expensive scans, duplicate intermediate state, and weak replay semantics.
- How to control complexity: fixed MVP rules, explicit windows, idempotent updates, per-rule metrics, and rollback flags.

06

Reliability Model And Sev-1 Runbook

Design for silent failure detection, not just service uptime




■ Incident Step	SEM / IC Focus	Technical Focus
Declare	Single IC, comms owner, customer impact clock	Freeze risky deploys
Scope	Tenant, region, event type, time window	Compare stage metrics
Mitigate	Choose customer-safe degraded mode	Failover, pause rules, restart consumers
Recover	Track backfill completion	Replay from durable offsets
Prevent	Postmortem and owners	Synthetic gaps, gates, runbooks

■ Operating principle: a stale green Health Center is worse than an explicit degraded state.

07 ■ Roadmap: MVP, GA, Platformization

Multi-quarter plan with customer-safety gates



■ MVP Gate

End-to-end synthetic, top 4 signals, SLO dashboards,

■ GA Gate

Scale, replay, correctness, shadow validation, incident

■ Migration Gate

Measured FP/FN rate and rollback path per tenant cohort.

■ Dependency Delay And Operational Drain

Absorb real-world disruption without pretending capacity is infinite


■ Gateway Delayed By Two Months

Freeze telemetry contract, build adapter boundary, generate synthetic/replayed streams, negotiate thin-slice routing for highest-value signals.

■ Legacy Offline Escalations +40%

Treat as customer trust risk; stabilize with a short-lived lane, not a full-team roadmap derailment.

■ 10-person team allocation during escalation



■ Exit Criteria

Legacy lane ends after top root cause patch, instrumentation, and support-volume trend review.

■ Escalation Rule


Increase allocation only for active Sev-1/Sev-2 impact across major tenants.

■ Learning Loop

Each false-offline root cause becomes a Health Center test, rule refinement, or migration guardrail.

■ Technical Conflict Resolution

Use evidence, decision criteria, and ownership to align senior engineers



◆ **Expected decision: stream processing for coalescing, constrained to MVP rules and strong operability controls.**

- Engineer A leads processing design and SLO instrumentation.
- Engineer B leads simplicity controls: bounded scope, failure-mode review, cost model, and rollback plan.

A 3x3 grid representing a 2D plane with axes 'Likelihood' (vertical) and 'Distance' (horizontal). The grid contains colored squares with letters: Top row (S, G), Middle row (N, Q, L), Bottom row (empty). A yellow square is at the bottom-left, and a green square is at the bottom-left of the yellow square.

Adapter contract, synthetic streams, thin-slice integration.

Two-engineer stabilization lane with exit criteria.

Coalescing, suppression, tenant rollout, support feedback

Narrow MVP rules, SLOs, runbooks, rollback flags.

Precomputed read models, bounded APIs, freshness

- More trustworthy agent-health visibility for customers and support.
- Reduced false-offline escalations and less legacy operational drag.
- Reusable Agent Platform foundation for future health and alerting use cases.

Presenter Notes Appendix

Detailed talking points, anticipated Q&A, and transition cues for each slide.

SLIDE 01 PRESENTER NOTES — Slide 1: Executive Framing

OPENING (60 seconds)

- Start by anchoring on the CUSTOMER problem, not the technology.

"Customers today cannot reliably tell whether their SentinelOne agents are healthy."

- Frame the assignment as a leadership test, not just architecture.

KEY TALKING POINTS

1. Four outcomes — Customer / Platform / Execution / Org — show you understand this is more than a feature build.
2. KPIs are concrete and measurable. Call out the <5-minute detection freshness as the hardest constraint that drives the architecture.
3. Leadership thesis: staged scope, evidence-based decisions, protect current customers. This separates you from candidates who only talk technology.

ANTICIPATED QUESTIONS

- "Why not start with all anomaly types?"
 - Explain ruthless MVP scoping; we prove correctness on 4 signals first.
- "How do you balance new-platform velocity with legacy maintenance?"
 - Preview the dedicated stabilization lane (Slide 8).

TRANSITION → "Let me show you how I scoped the MVP and set SLOs before discussing architecture choices."

SLIDE 02 PRESENTER NOTES — Slide 2: Scope, SLOs & Guardrails

PURPOSE

Show the audience you scope BEFORE designing. Most SEM candidates jump to architecture; you start with requirements discipline.

KEY TALKING POINTS

1. MVP limits to 4 anomaly types — high signal, customer-actionable.
Everything else is explicitly deferred to reduce delivery risk.
2. SLO table: call out "silent gap detection" — this is the most dangerous failure mode in a security product (stale green dashboards).
3. GA scope adds alert lifecycle, suppression, and shadow migration.
Shadow validation is critical: we do NOT flip source-of-truth until we can quantify false-positive/negative rates.
4. Guardrail: prevents scope creep. Every feature request during MVP must pass the "actionability or de-risk" test.

DATA POINTS TO MENTION

- Billions of events/day at millions of agents.
- Multi-tenant isolation with fair-usage controls.
- At-least-once delivery with idempotent processing.

ANTICIPATED QUESTIONS

- "Why not add notifications in MVP?"
→ Notifications add blast-radius risk; we need correctness first.
- "How do you measure the 5-minute freshness SLO?"
→ Synthetic events flowing end-to-end (preview Slide 6).

TRANSITION → "With scope locked, let me walk you through the reference architecture that satisfies these SLOs."

PURPOSE

Present a cloud-neutral, replayable architecture. Emphasize the SEPARATION OF CONCERNS (hot / read / cold) — this is what makes the system scalable and operable.

WALK-THROUGH (left to right, ~90 seconds)

1. Agents → Ingestion Gateway: authentication, rate-limiting, routing.
Gateway team owns this; we consume via a versioned telemetry contract.
2. Durable Event Bus: partitioned by tenant + agent ID.
At-least-once delivery; replay is the key capability.
3. Stream Processing → Rules + Coalescing: real-time anomaly detection.
This is where the 5-min SLO is enforced.
4. Alert State Store: write-optimized, tenant/agent-keyed.
5. Read Model + Search: pre-aggregated views for console queries.
6. APIs + Console: bounded queries, pagination, freshness metadata.

COLD PATH

- Data Lake stores every event for audit, replay, and rule tuning.
- Lifecycle Store tracks alert history for compliance.

TECHNOLOGY CHOICES

- Deliberately cloud-neutral labels. In the interview, mention:
Kafka or PubSub, Flink or Kafka Streams, DynamoDB or Cassandra, OpenSearch or Elasticsearch, S3/GCS for the data lake.
- Do NOT over-commit to specific vendors; show flexibility.

DESIGN PRINCIPLES

- Write path never shares resources with read path.
- Every stage emits metrics for pipeline-health monitoring (Slide 6).
- Replay-from-offset is the primary recovery mechanism.

ANTICIPATED QUESTIONS

- "Why not a single database?"
→ Different access patterns require different stores; forcing one store to serve writes + queries + analytics creates contention.
- "How do you handle schema evolution?"
→ Versioned telemetry contract with backward-compat validation.

TRANSITION → "Let me zoom into the read path and how we hit the 200ms

latency target."

SLIDE 04 PRESENTER NOTES — Slide 4: Read Path & Latency Budget

PURPOSE

Demonstrate that the 200ms p95 target is achievable through design, not hope. Show you've thought through every millisecond of the request path.

WALK-THROUGH

1. Console → APIs → Read Model → Search → Alert State.

Data flows RIGHT to LEFT in the query path.

2. Latency budget adds up to ~150–220 ms. We have headroom under p95.

3. Key insight: the UI queries PRECOMPUTED VIEWS, not raw events.

IMPORTANT DESIGN CHOICES

- Graceful Degradation: if search index lags, the API returns current alert state with freshness metadata. The console shows "data may be delayed" instead of stale green. This is critical for a security product.

- Query Discipline: bounded queries prevent runaway scans. No "select * from events where tenant=X" from the dashboard.

- Correctness: at-least-once + idempotent transitions means we can safely replay without creating duplicate alerts.

LATENCY BUDGET DEEP-DIVE (if asked)

- Auth/gateway: JWT validation, tenant-scoped claims, rate shaping.
- Orchestration: single API call to the read model, no service fanout.
- Read store: OpenSearch query with pre-aggregated alert-count views.
- Serialization: protobuf or slim JSON, paginated to max 100 items.

ANTICIPATED QUESTIONS

- "What if a tenant has 100K agents?"

→ Pagination + server-side aggregation. Dashboard shows top-N with drill-down, not a flat list.

- "How do you cache?"

→ Tenant-level dashboard summaries with TTL-based invalidation. Per-agent pages are not cached.

TRANSITION → "Now let me show you the hardest technical problem: alert coalescing within the 5-minute SLO."

PURPOSE

This is the HARDEST TECHNICAL PROBLEM in the design. Show that you understand the trade-off and can make an evidence-based decision.

EXAMPLE WALK-THROUGH (point to timeline)

- 50 anti-tamper events from one agent over 10 minutes.
- Without coalescing, the dashboard shows 50 separate alerts → noise.
- With coalescing: key by tenant+agent → tumbling window → state machine

emits ONE alert with evidence count.

STATE MACHINE TRANSITIONS

Open → Update Evidence → Suppress Duplicate → Resolve (healthy signal)

Each transition is idempotent using (tenant, agent, rule, window) key.

WHY STREAM PROCESSING OVER CRON

1. Freshness: the 5-min SLO cannot be reliably met with periodic batch.
2. Scale: scanning billions of events in cron is expensive.
3. Replay: stream offsets give free replay; cron needs custom backfill.
4. Intermediate state: cron creates half-processed records.

COMPLEXITY CONTROLS (critical for credibility)

- Only 4 rules in MVP — no generic rule engine.
- Explicit tumbling windows (configurable per rule).
- Idempotent state updates.
- Per-rule metrics: processing lag, emit rate, error rate.
- Rollback flags: disable individual rules without redeploying.

THIS IS THE ENGINEER A vs ENGINEER B DEBATE

- Preview that you expect stream processing to win, but you'll run a structured decision process (Slide 9).

ANTICIPATED QUESTIONS

- "What if a rule is buggy?"
 - Feature flag per rule. Disable instantly. Roll back to last-known-good rule version.
- "How do you test coalescing logic?"
 - Unit tests with synthetic event sequences. Integration tests with replayed production data (scrubbed).

TRANSITION → "How do we know all of this is working correctly in

production? Let me walk through the reliability model."

SLIDE 06 PRESENTER NOTES — Slide 6: Reliability & Sev-1 Runbook

PURPOSE

This slide is CRITICAL. It demonstrates operational maturity — the #1 differentiator for a Senior EM candidate.

OBSERVABILITY PIPELINE (top diagram)

Walk through each stage LEFT to RIGHT:

1. Ingress rate: events accepted by the gateway per second.
2. Bus lag: consumer offset lag on the durable event bus.
3. Processor lag: processing latency vs event time.
4. Store writes: write throughput to alert state store.
5. Index freshness: lag between alert-store writes and search-index availability.
6. Synthetic probe: end-to-end synthetic events per tenant cohort.
7. API p95: dashboard response time.

KEY INSIGHT: "A Kubernetes pod being UP is not proof that the pipeline is processing events correctly."

SEV-1 RUNBOOK — WALK THROUGH EACH ROW

1. DECLARE: Single incident commander. Start customer-impact clock. Freeze risky deployments.
2. SCOPE: Determine blast radius — which tenants, regions, event types, time window. Compare metrics across stages.
3. MITIGATE: Choose a CUSTOMER-SAFE degraded mode. Pause bad rules. Restart consumers. Failover to standby.
4. RECOVER: Replay from durable offsets. Backfill missing alerts. Track completion.
5. PREVENT: Blameless postmortem. Add synthetic gaps. Improve gates.

OPERATING PRINCIPLE (bottom, in red)

"A stale green Health Center is worse than an explicit degraded state."

This is the single most important sentence in the deck. It shows you understand security-product operations.

ANTICIPATED QUESTIONS

- "How long does replay take?"
→ Depends on gap duration. For a 1-hour gap with billions of events, expect 20-30 minutes with parallelized consumers.
- "Who gets paged?"

→ On-call SRE + Health Center tech lead + SEM for customer comms.

TRANSITION → "Let me show you the multi-quarter roadmap with delivery gates."

PURPOSE

Show a credible, multi-quarter execution plan with EXPLICIT GATES that prevent premature rollout.

SWIM-LANE WALK-THROUGH (~60 seconds)

- Telemetry + Processing: contract → adapt → top 4 rules (Q1) → scale/replay/idem (Q2) → reusable framework (Q3) → broader catalog (Q4).
- Console + APIs: dashboard/alerts/flags (Q1) → lifecycle/search (Q2) → trends/reports/notifs (Q3) → legacy UI consolidation (Q4).
- Operations: SLOs/synthetic/runbooks from DAY ONE (Q1) → game days/canaries (Q2) → cost model/onboarding (Q3) → decom runbooks (Q4).
- Legacy Offline: stabilize false-offline (Q1) → shadow compare (Q2) → tenant migration (Q3) → retire safe paths (Q4).

GATE DEFINITIONS (bottom cards)

1. MVP Gate: cannot GA without end-to-end synthetic events, SLO dashboards, and controlled allowlist rollout.
2. GA Gate: cannot go broad without scale testing, replay validation, shadow comparison, and incident readiness.
3. Migration Gate: cannot retire legacy without measured FP/FN rates and a per-tenant rollback path.

KEY MESSAGE: "Operations is a first-class swim lane, not an afterthought. We ship SLOs and runbooks before we ship features."

ANTICIPATED QUESTIONS

- "Can you compress Q1 and Q2?"
→ Possibly, if gateway dependency resolves early. But I'd rather ship a bulletproof MVP than a rushed GA.
- "When is legacy fully retired?"
→ Q4 for safe paths; some long-tail tenants may extend into Q5.

Explicit tenant-by-tenant migration with rollback.

TRANSITION → "Let me address the two biggest disruptions: the gateway delay and the legacy escalation spike."

PURPOSE

This is where you demonstrate EXECUTION LEADERSHIP. Two real-world disruptions that most candidates handle poorly:

1. A dependency team is delayed.
2. A legacy system is burning.

GATEWAY DELAY RESPONSE

- Do NOT idle the team.
- Freeze a versioned telemetry contract.
- Build an adapter boundary against that contract.
- Generate synthetic + replayed streams.
- Unblock: processors, schemas, APIs, UI, QA, observability, load tests.
- Negotiate thin-slice integration: route 1-2 high-value signals first.
- Only dependency-bound scope shifts; everything else continues.

LEGACY ESCALATION RESPONSE

- Treat as customer trust risk, NOT background maintenance.
- Create a SHORT-LIVED stabilization lane: 2 engineers + QA support.
- Keep 6 engineers on Health Center MVP. 1 tech lead coordinates.
- Mission: instrument → root-cause → patch → reduce support volume.
- Exit criteria are explicit: lane closes after patch + instrumentation +

support-volume trend review.

TEAM ALLOCATION BAR (visual)

- Walk through: 6 build / 2 legacy / 1 QA / 1 coordination = 10 total.
- Emphasize: this is a TEMPORARY allocation, not a permanent split.
- Escalation rule: only increase legacy if Sev-1/Sev-2 across major tenants.

LEARNING LOOP

- Every false-offline root cause feeds back into Health Center:

becomes a test case, a rule refinement, or a migration guardrail.

- This turns the legacy pain into strategic value.

ANTICIPATED QUESTIONS

- "What if the gateway delay extends to 4 months?"

→ We can ship MVP entirely on synthetic + existing heartbeat streams.

Production gateway integration becomes a GA gate.

- "What if escalations consume the whole team?"

→ The escalation rule is explicit. SEM protects roadmap unless

active Sev-1/Sev-2 across major tenants.

TRANSITION → "Now let me address the people challenge: resolving the technical conflict between the two senior engineers."

PURPOSE

This is the PEOPLE LEADERSHIP test. The interviewer wants to see that you can lead senior engineers through disagreement WITHOUT:

- Pulling rank ("just do it my way")
- Avoiding the decision ("let's try both")
- Making it personal ("A is right, B is wrong")

ORG CHART VISUAL (top)

- Engineer A: advocates stream processing (Kafka Streams / Flink).
- Engineer B: advocates DB writes + cron aggregation.
- SEM (you): turns disagreement into a decision system.

4-STEP PROCESS

1. CRITERIA: Define decision criteria BEFORE debating solutions.

Freshness, scale, cost, operability, delivery risk.

2. SPIKE: One week with realistic data. Not a toy prototype.

Realistic cardinality (millions of agents), duplicate-event rates, and failure cases (what happens when a processor crashes?).

3. DECISION: Architecture Decision Record (ADR).

Document: assumptions, constraints, decision, trade-offs, rollback trigger, revisit conditions.

4. COMMIT: Both engineers own part of the winning design.

Engineer A leads the processing layer.

Engineer B leads simplicity controls and rollback planning.

This turns disagreement into a STRONGER DESIGN.

WHY STREAM PROCESSING WINS

- 5-minute freshness SLO at billions of events/day.
- Stream-time windows naturally model coalescing.
- Durable offsets enable replay.
- Cron risks: late detection, expensive scans, duplicate state.

BUT with constraints:

- Only 4 rules in MVP.
- No generic rule engine.
- Feature flags per rule.
- Strong operability: metrics, dashboards, runbooks, rollback.

ANTICIPATED QUESTIONS

- "What if they still disagree after the spike?"

→ The SEM makes the call. The ADR documents it. Both engineers sign off on the document, not just the decision.

- "Have you done this before?"

→ [Prepare a real example from your experience]

TRANSITION → "Let me close with the major risks, mitigations, and the outcomes we expect to deliver."

SLIDE 10 PRESENTER NOTES — Slide 10: Risks, Mitigations & Outcomes

PURPOSE

Close strong. This slide should leave the interviewer confident that you can OPERATE, not just design.

RISK MATRIX WALK-THROUGH

Point to each dot and explain:

G (Gateway Delay) — HIGH likelihood, HIGH impact.

Mitigation: adapter contract, synthetic streams, thin-slice integration.

Why it's manageable: we decouple downstream development from the gateway timeline.

S (Stream Complexity) — MEDIUM likelihood, HIGH impact.

Mitigation: narrow MVP rules, SLOs, runbooks, rollback flags.

Why it's manageable: complexity is bounded by scope control.

L (Legacy Drain) — HIGH likelihood, MEDIUM impact.

Mitigation: dedicated stabilization lane with explicit exit criteria.

Why it's manageable: short-lived, bounded allocation.

Q (Query Latency) — MEDIUM likelihood, MEDIUM impact.

Mitigation: precomputed read models, bounded APIs, freshness metadata.

Why it's manageable: standard read-path optimization.

N (Noise / Alert Fatigue) — LOW likelihood, MEDIUM impact.

Mitigation: coalescing, suppression, tenant rollout, feedback loop.

Why it's manageable: tenant-by-tenant rollout catches noise early.

EXPECTED OUTCOMES (bottom)

1. More trustworthy agent-health visibility.
2. Reduced false-offline escalations → less support burden.
3. Reusable Agent Platform foundation → future leverage.

CLOSING STATEMENT (memorize this)

"The Health Center program succeeds if it improves customer trust while creating a durable Agent Platform foundation. The architecture handles hyperscale telemetry, but the bigger test is execution judgment: narrow MVP scope, protect customer-facing reliability, absorb dependency delays, resolve technical conflict with evidence, and keep the team engaged through operational pressure."

ANTICIPATED QUESTIONS

- "What would you cut if you had only one quarter?"

→ Ship top 2 signals (offline + anti-tamper) with synthetic SLOs on allowlisted tenants. No lifecycle, no migration, no search.

- "How do you measure success?"

→ 50% reduction in false-offline escalations within 2 quarters.

SLO dashboards showing 99%+ freshness compliance.

- "What's the biggest lesson from a past project?"

→ [Prepare a real example]