

Strategic Architectural & Operational Blueprint:

Singularity Health Center

1. Executive Framing and Platform Vision

Singularity Health Center (SHC) is not merely a monitoring tool; it is a high-scale streaming state platform engineered to maintain foundational trust in a multi-tenant cybersecurity ecosystem. In the SentinelOne environment, the health of the agent is the bedrock of our security posture. We have strategically decoupled health telemetry from primary security detection pipelines to ensure operational visibility is never compromised during "intense security events," where security pipelines may be throttled or saturated. By isolating agent health processing, we guarantee a dedicated, high-availability path for health telemetry that is immune to spikes in malicious activity.

Our primary engineering mission is to detect agent drift, connectivity loss, and anti-tamper breaches with authoritative precision. Conversely, SHC explicitly excludes malicious activity detection, arbitrary customer analytics, or direct endpoint remediation in its initial phase. These boundaries prevent the "scope creep" that typically destabilizes high-volume ingestion systems and allow our team to optimize for the unique throughput and reliability constraints of global telemetry.

Capacity Assumptions & Design Implications

The following parameters define the scale required to maintain a trustworthy posture across the global fleet.

Dimension	Example Assumption	Design Implication
Managed Agents	10 million	State store must handle tens of millions of "hot keys" for per-agent health status.
Telemetry Events	5 billion / day	Average ingestion of ~58k events/sec necessitates horizontal scaling and durable buffering.
Peak Throughput	10x multiplier	Must support ~580k events/sec regional burst capacity to handle incident-driven spikes.
Data Volume	25–100 TB / day	High ingest envelope requires compressed raw storage and lean, versioned state-update logic.

Detection Latency	P95 < 2–5 minutes	Direct signal detection (seconds/mins) vs. Heartbeat anomalies (Threshold + 2 mins).
Alert Fanout	1–5% of agents/day	Large-scale incidents require mandatory deduplication and incident grouping/coalescing.

These high-level requirements necessitate a rigorous evaluation of technical trade-offs, shifting our focus from standard microservice patterns to the specialized primitives of a global-scale streaming architecture.

2. Deep Technical Trade-offs and Architectural Strategy

At a scale of 100 TB per day, traditional architectural "best practices" often conflict. We cannot optimize for every variable simultaneously; we must prioritize technical compromises that balance latency, cost, and correctness. This requires a deliberate departure from general-purpose solutions in favor of specialized streaming strategies.

Partitioning for Multi-Tenant Isolation

Our primary defense against "Hot Tenants"—enterprise customers who may generate 30% or more of total volume—is our partition key strategy. We utilize `tenant_id + hash(agent_id)` for stream partitioning. This preserves per-agent event ordering while effectively sharding large tenants across the processing cluster, preventing a single noisy customer from monopolizing resources or inducing lag in others.

Streaming vs. Batch Evaluation

We have implemented a streaming-first approach for detection to satisfy our 5-minute Service Level Objective (SLO). While batch processing is simpler, it introduces unacceptable "dead time" that violates the needs of a real-time security platform. The data lake remains a secondary, critical asset for historical replay, auditability, and backtesting new health rules against real-world data without impacting the live production stream.

Hot State Management: KV Stores vs. Search Engines

We have opted for Key-Value (KV) stores (e.g., DynamoDB/Cassandra) over Search Engines (e.g., Elasticsearch) for maintaining the authoritative "Hot State."

- **Point-Read Efficiency:** KV stores provide predictable, sub-millisecond point-read and write performance, critical for updating millions of agent records every few minutes.
- **Write Amplification:** Search engines suffer from massive write amplification and indexing overhead at 100 TB/day, making them cost-prohibitive for authoritative state.
- **Asynchronous Investigation:** In this architecture, search engines are relegated to asynchronous investigation and historical drill-down rather than serving as the source of truth for online/offline status.

The "Effectively-Once" Processing Model

We have discarded end-to-end "Exactly-Once" processing as a theoretical ideal that is too fragile for this scale. We instead utilize an Effectively-Once model built on three pillars:

- 1. Idempotency Keys: Unique fingerprints ensure that retired events do not trigger duplicate alerts.
- 2. Monotonic State Versions: We use versioning to reject stale state transitions. The state updater rejects any update with a version lower than the current state, ensuring correctness regardless of event-time arrival order.
- 3. Deduplication Windows: Events are buffered in short time windows to coalesce repeated signals before they reach the state store.

This architectural idealism must now meet the concrete operational realities of managing our legacy infrastructure and the "drain" of technical debt.

3. Operational Resilience and Managing Legacy Drain

Building a next-generation platform while supporting a fragile legacy system is a leadership challenge in stabilization. Our legacy "heartbeat" microservice is currently causing significant customer friction, manifested as a 40% increase in false-positive "offline" alerts. We are establishing a stabilization lane to protect customer trust while maintaining velocity on the new platform's critical path.

Managing the Legacy Offline Status Escalation

To resolve the current instability, we have deployed a "Tiger Team" with a focused remediation mandate:

- 1. Diagnosis: Identifying whether the 40% spike stems from aggressive heartbeat timeouts, metadata mismatches, or UI cache staleness.
- 2. Acute Patching: Implementing a "kill switch" for the most aggressive alerting rules and adjusting heartbeat thresholds to provide a temporary buffer.
- 3. Backporting Logic: Backporting the freshness and monotonic state-versioning logic from SHC into the legacy service to stabilize it until decommission.

Sev-1 Scenario: Silent Processing Stall

The most dangerous failure is a silent stall, where the UI shows a "healthy" fleet because the pipeline has stopped updating. We utilize the following diagnostic decision tree:

Symptom	Likely Cause	Action
Rising Consumer Lag	Downstream store throttling	Scale KV store capacity; check for hot partitions.
Synthetic Canary Failure	Logic regression or parser crash	Roll back the most recent deployment; inspect DLQ for "poison events."

Stale Read-Model Timestamps	API/Cache layer failure	Invalidate Redis cache; force rebuild of precomputed read models.
Ingest Rate vs. Process Rate Mismatch	Network partition or Gateway issue	Correlate with Ingestion Gateway metrics; surface freshness warning in UI.

Risk Register

Risk	Impact	Architectural Mitigation
Alert Storms	Customer fatigue/noise	Grouped incident detection and stable idempotency keys.
Hot Tenants	Partition imbalance/lag	Partitioning by <code>tenant_id + hash(agent_id)</code> and dedicated lanes.
Metadata Lag	Inaccurate drift alerts	Treating metadata freshness as a confidence score in the rule engine.
False Positives	Loss of customer trust	"Shadow Mode" validation and platform-incident suppression logic.

4. People Management and Technical Leadership

At the Senior Director level, resolving high-stakes technical disputes is about providing a decision framework rather than picking a "winner." I prioritize evidence-based consensus to ensure the team remains unified.

The Alert Coalescing Conflict

A dispute recently arose between Engineer A (Streaming) and Engineer B (Async/Cron) regarding event grouping.

- Engineer A (Streaming): Proposes coalescing alerts within the stream (Flink/Kafka).
 - *Strengths*: Low latency; meets 5-minute SLO; scales with partition keys.
 - *Risks*: Increased infrastructure complexity.
- Engineer B (Async/Cron): Proposes writing raw events and using a cron job to group them later.
 - *Strengths*: Simpler initial implementation.
 - *Risks*: Fails the 5-minute SLO; induces massive database scan costs at scale.

Mediation and Resolution via ADR

To resolve this, I mandated a time-boxed technical spike to prototype both approaches. Following a review of the decision matrix, we have formally adopted stream-time coalescing via an Architecture Decision Record (ADR). This decision was justified by our non-negotiable 5-minute detection SLO and the need to prevent notification spam before it reaches the customer-facing database.

Morale and Growth Strategy during Dependency Delays

With the two-month delay of the Ingestion Gateway, morale is at risk. I have framed this delay not as a "wait," but as a Growth Opportunity. The team is currently owning high-impact deliverables that ensure "plug-and-play" readiness:

- ADR & Load Test Harnesses: Developing the rigorous testing frameworks needed to survive 10x burst events.
- Simulators: Building high-fidelity telemetry generators to stress-test the normalization logic.
- UI Mocks & API Contracts: Finalizing the React console views and mocking read models to ensure frontend-backend alignment.

By focusing on these areas, the team maintains technical agency and ensures a rapid transition to GA once the dependency is resolved.

5. Roadmap, Success Metrics, and GA Readiness

Our rollout follows a phased approach to build customer trust incrementally through "Shadow Mode" and "Private Previews."

Milestone Plan

Phase	Milestone	Core Deliverables	Exit Criteria
M0	Discovery	Event taxonomy, capacity model.	Reviewed design & sized backlog.
M1	Foundation	Ingest gateway contracts, state service.	Process synthetic telemetry at peak.
M2	Rules/Alerts	Versioned rules, deduplication logic.	Replay can reproduce outputs for a tenant/time range.
M3	Preview	Dashboard UI, lifecycle (Ack/Mute).	Preview customers use the console unassisted.

M4	GA	Notifications, canary rollouts.	End-to-end SLOs met under full load.
----	----	---------------------------------	--------------------------------------

Key Performance Indicators (KPIs)

Category	Metric	Target / Goal
Product Outcomes	Alert Acknowledgement Rate	A high rate indicates actionable, trusted alerts.
Engineering SLOs	P95 Detection Latency	Under 5 minutes for all critical health signals.
Customer Experience	API Latency	< 200ms for all summary and list views.

GA Readiness Checklist

Before General Availability, the platform must pass these final gates:

- Sev-1 Game Days: Simulating regional outages and silent processing stalls to validate runbooks.
- Rule Suppression: Validation of incident-suppression logic during simulated platform outages.
- Decommission Path: A finalized migration of all legacy "online/offline" traffic to the new SHC read models, including a dual-read validation period to ensure data parity.

This strategic plan transforms raw, high-volume telemetry into a high-value, trustworthy operational asset, ensuring Singularity Health Center becomes the definitive source of truth for the health of the SentinelOne fleet.